

# EDA on Haberman's Survival Data Set

(Survival of patients who had undergone surgery for breast cancer)

## About the Dataset

### Haberman's data set

The dataset contains cases from a study that was conducted between 1958 and 1970 at the University of Chicago's Billings Hospital on the survival of patients who had undergone surgery for breast cancer.

Reference : <https://www.kaggle.com/gilsousa/habermans-survival-data-set>  
(<https://www.kaggle.com/gilsousa/habermans-survival-data-set>)

### Feature/Variable/Input-variable/Independent-varibale

There are 4 attributes: 'Age', 'Op\_Year', 'axil\_nodes', 'Surv\_status'

- Feature Information:
- 'Age'-----Age of patient at time of operation (numerical)
- 'Op\_Year'-----Patient's year of operation (year - 1900, numerical)
- 'axil\_nodes'--Number of positive axillary nodes detected (numerical)

### Label/Dependent-variable

- Survival status (class attribute)
  - 1 = the patient survived 5 years or longer
  - 2 = the patient died within 5 years

```
In [18]: import pandas as pd
import seaborn as sbn
import matplotlib.pyplot as plt
import numpy as np
```

```
In [50]: import warnings
warnings.filterwarnings("ignore")
```

```
In [4]: #Load haberman.csv into a pandas DataFrame.
df = pd.read_csv("haberman.csv")
df.head()
```

Out[4]:

	30	64	1	1.1
0	30	62	3	1
1	30	65	0	1
2	31	59	2	1
3	31	65	4	1
4	33	58	10	1

```
In [6]: #Rename the column name
df = df.rename(columns={'30': 'Age', '64': 'Op_Year', '1': 'axil_nodes', '1.1': 'Surv_status'})
df.columns
```

Out[6]: Index(['Age', 'Op\_Year', 'axil\_nodes', 'Surv\_status'], dtype='object')

```
In [8]: #Data-point
print(df.shape)

print("=="*10)
print("Data points per each class: ")
df['Surv_status'].value_counts()
```

```
(305, 4)
=====
Data points per each class:
```

```
Out[8]: 1    224
        2     81
        Name: Surv_status, dtype: int64
```

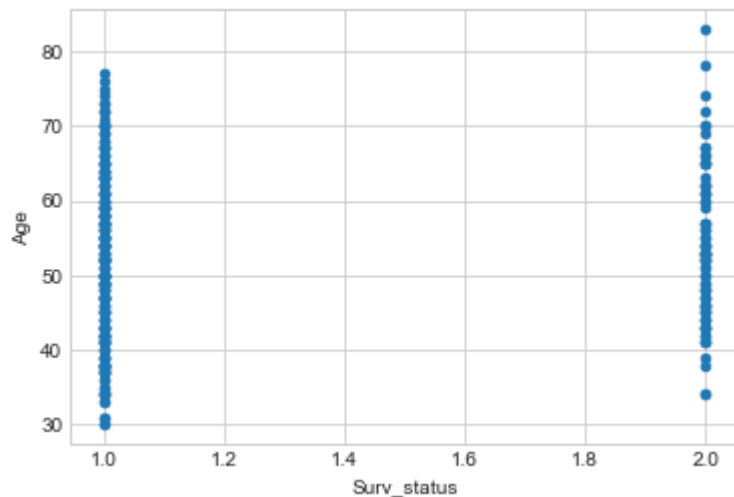
## Observation:

- There are two classes 1 and 2, class 1 denotes patients who survived 5 years or longer after the operation and class 2 denotes to patients who died within 5 years of operation.
- class 1 contains 224 values and class 2 contain 81 values
- It is being clear that 224 patients survived 5 years or longer after the operation and 81 patients died within 5 years of operation.


# Bi-variate analysis

## 1) 2-D Scatter Plot


```
In [21]: #Age  
df.plot(kind="scatter", x="Surv_status", y="Age")  
plt.show()
```



```
In [22]: #Op_Year  
df.plot(kind="scatter", x="Surv_status", y="Op_Year")  
plt.show()
```



```
In [23]: #axil_nodes
df.plot(kind="scatter", x="Surv_status", y="axil_nodes")
plt.show()
```




## Observation

- This can't make much sense out of it.
- we can color the points by their Feature/Survival Status.

## 2-D Scatter plot with color-coding

```
In [17]: sbn.set_style("whitegrid");
sbn.FacetGrid(df, hue="Surv_status", size=4) \
    .map(plt.scatter, "Age", "Op_Year") \
    .add_legend();
plt.show();
```



## Observation

- We can plot 2-D Scatter plot with color-coding for each Feature/Survival Status
- It's a little bit difficult to identify the plots one by one.

## 2) Pair-plot

- We can draw multiple 2-D scatter plots for each combination of features.

```
In [14]: sbn.pairplot(df, hue="Surv_status", size=3)
plt.show()
```




## Observations

- axil\_nodes and Surv\_status are the most useful features to identify Survival status of the patients.
- We can use "line: between both" and "if-else" conditions to build a simple model to classify patients status.


## Histogram

- 1-D scatter plot

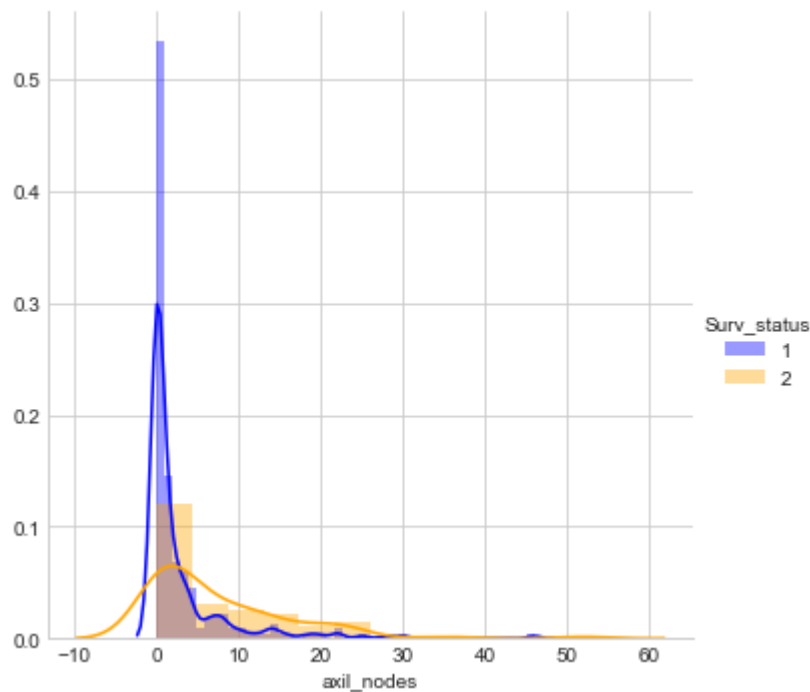
```
In [59]: #Age
sbn.set_style("whitegrid")
pal = ["blue","orange"]
sbn.FacetGrid(df, hue="Surv_status", palette= pal,size = 5)\
    .map(sbn.distplot, "Age")\
    .add_legend()
plt.show()
```



```
In [61]: #Op_Year
sbn.set_style("whitegrid")
sbn.FacetGrid(df, hue= "Surv_status", palette= pal, size= 5)\
    .map(sbn.distplot, "Op_Year")\
    .add_legend()
plt.show()
```



```
In [62]: #axil_nodes
sbn.set_style("whitegrid")
sbn.FacetGrid(df, hue= "Surv_status", palette= pal, size= 5)\
    .map(sbn.distplot, "axil_nodes")\
    .add_legend()
plt.show()
```



## Univariate Analysis using PDF

- one variable analysis

```
In [10]: df_survived = df.loc[df["Surv_status"] == 1]
df_died = df.loc[df["Surv_status"] == 2]
```



```

In [66]: #Age
counts, bin_edges = np.histogram(df_survived['Age'], bins=10,
                                density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)


counts, bin_edges = np.histogram(df_died['Age'], bins=10,
                                density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)

plt.show()

[0.04910714 0.10714286 0.125      0.09375    0.16517857 0.16517857
 0.09375    0.11160714 0.0625     0.02678571]
[30.  34.7 39.4 44.1 48.8 53.5 58.2 62.9 67.6 72.3 77. ]
[0.03703704 0.12345679 0.19753086 0.19753086 0.13580247 0.12345679
 0.09876543 0.04938272 0.02469136 0.01234568]
[34.  38.9 43.8 48.7 53.6 58.5 63.4 68.3 73.2 78.1 83. ]

```



## Observation:

- Univariate analysis based on Age gives a slightly different plot for both classes with some overlap of data points.
- it is difficult to set a threshold to distinguish the two classes.
- Age is not a very important feature for determining the survival of a patient.

```

In [67]: #Op_Year
counts, bin_edges = np.histogram(df_survived['Op_Year'], bins=10,
                                density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)

counts, bin_edges = np.histogram(df_died['Op_Year'], bins=10,
                                density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)


plt.show()

```

```

[0.1875      0.10714286 0.10267857 0.07142857 0.09821429 0.09821429
 0.06696429 0.09821429 0.09375    0.07589286]
[58.  59.1 60.2 61.3 62.4 63.5 64.6 65.7 66.8 67.9 69. ]
[0.25925926 0.04938272 0.03703704 0.08641975 0.09876543 0.09876543
 0.16049383 0.07407407 0.04938272 0.08641975]
[58.  59.1 60.2 61.3 62.4 63.5 64.6 65.7 66.8 67.9 69. ]

```



## Observation:

- Univariate analysis based on Year\_of\_operation(Op\_Year) gives similar plots for both classes with a huge overlap rate of data points.
- Year\_of\_operation(Op\_Year) is not an important feature for determining the survival of a patient.
- it's also difficult to set a threshold to distinguish the two classes.

```

In [65]: #axil_nodes
counts, bin_edges = np.histogram(df_survived['axil_nodes'], bins=10,
                                density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)

counts, bin_edges = np.histogram(df_died['axil_nodes'], bins=10,
                                density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)


plt.show()

```

```

[0.83482143 0.08035714 0.02232143 0.02678571 0.01785714 0.00446429
 0.00892857 0.          0.          0.00446429]
[ 0.   4.6  9.2 13.8 18.4 23.   27.6 32.2 36.8 41.4 46. ]
[0.56790123 0.14814815 0.13580247 0.04938272 0.07407407 0.
 0.01234568 0.          0.          0.01234568]
[ 0.   5.2 10.4 15.6 20.8 26.   31.2 36.4 41.6 46.8 52. ]

```



## Observation:

- Univariate analysis based on axil\_nodes gives less overlap of data points as compared to other features.
- It's also difficult to set a threshold for axillary nodes which will differentiate both classes of patients.
- axil\_nodes can be an important feature for determining the survival of a patient.

## Mean, Variance and Std-dev

In [83]: `print(df_survived.describe())`

	Age	Op_Year	axil_nodes	Surv_status
count	224.000000	224.000000	224.000000	224.0
mean	52.116071	62.857143	2.799107	1.0
std	10.937446	3.229231	5.882237	0.0
min	30.000000	58.000000	0.000000	1.0
25%	43.000000	60.000000	0.000000	1.0
50%	52.000000	63.000000	0.000000	1.0
75%	60.000000	66.000000	3.000000	1.0
max	77.000000	69.000000	46.000000	1.0

In [86]: *#Median is equivalent to Mean*  
`print("Median from df_survived")`  
`print(np.median(df_survived["Age"]))`  
`print(np.median(df_survived["Op_Year"]))`  
`print(np.median(df_survived["axil_nodes"]))`  
`print("=="*30)`  
`print("Median from df_died")`  
`print(np.median(df_died["Age"]))`  
`print(np.median(df_died["Op_Year"]))`  
`print(np.median(df_died["axil_nodes"]))`

Median from df\_survived

52.0

63.0

0.0

=====

Median from df\_died

53.0

63.0

4.0

In [82]: `print("\n90th Percentiles:")`  
`print(np.percentile(df_survived["Age"],90))`  
`print(np.percentile(df_survived["Op_Year"],90))`  
`print(np.percentile(df_survived["axil_nodes"],90))`  
`print("=="*30)`  
`print(np.percentile(df_died["Age"],90))`  
`print(np.percentile(df_died["Op_Year"],90))`  
`print(np.percentile(df_died["axil_nodes"],90))`

90th Percentiles:

67.0

67.0

8.0

=====

67.0

67.0

20.0

```
In [92]: #MAD is equivalent to std-dev
print ("\nMedian Absolute Deviation")
from statsmodels import robust
print(robust.mad(df_survived["Age"]))
print(robust.mad(df_survived["Op_Year"]))
print(robust.mad(df_survived["axil_nodes"]))
print("=="*30)
print(robust.mad(df_died["Age"]))
print(robust.mad(df_died["Op_Year"]))
print(robust.mad(df_died["axil_nodes"]))
```

Median Absolute Deviation

13.343419966550417

4.447806655516806

0.0

=====


11.860817748044816

4.447806655516806


5.930408874022408

## Box plot and Whiskers


```
In [94]: sbn.boxplot(x="Surv_status", y="Age", data=df)
plt.show()
```



```
In [95]: sbn.boxplot(x="Surv_status", y="Op_Year", data=df)
plt.show()
```




```
In [96]: sbn.boxplot(x="Surv_status", y="axil_nodes", data=df)
plt.show()
```




## Violin plots

Boxplot + PDF


```
In [97]: sbn.violinplot(x= "Surv_status", y ="Age", data=df, size= 8)  
plt.show()
```



```
In [98]: sbn.violinplot(x= "Surv_status", y ="Op_Year", data=df, size= 8)  
plt.show()
```




```
In [99]: sbn.violinplot(x= "Surv_status", y ="axil_nodes", data=df, size= 8)  
plt.show()
```



## Multivariate probability density, contour plot.

```
In [104]: sbn.jointplot(x= "Surv_status", y="Age", data= df, kind= "kde")  
plt.plot()
```


Out[104]: []






```
In [103]: sbn.jointplot(x= "Surv_status", y="Op_Year", data= df, kind= "kde")  
plt.plot()
```

```
Out[103]: []
```



```
In [105]: sbn.jointplot(x= "Surv_status", y="axil_nodes", data= df, kind= "kde")  
plt.plot()
```

```
Out[105]: []
```



## Conclusions:

- Patient's age(Age) and year of operation(Op\_Year) only can not define the survival status that, the patient survived 5 years or longer or died within 5 years. it is difficult to predict.
- axillary(axil\_nodes) node plays important role to perform EDA on Haberman's Data Set.